

1

# Lab 1 – Hadoop & MapReduce (Word Count)

## 1 Create Docker Network

We create a bridge network to allow all Hadoop containers to communicate.

```
docker network create --driver bridge hadoop
```

```
docker network ls → To check available networks
```

## 2 Start Hadoop Master Container

Start the master container and map its Hadoop UI ports to host machine ports.

```
docker run -itd --network hadoop `  
  --hostname hadoop-master `  
  --name hadoop-master `  
  -p 8090:50070 `  
  -p 8091:8088 `  
  -p 16011:16010 `  
  liliafxi/spark-hadoop:hv-2.7.2
```

- **Ports mapping:**

- 50070 → 8090 : NameNode Web UI
- 8088 → 8089 : YARN ResourceManager Web UI
- 16010 → 16011 : JobHistory Server

## 3 Start Hadoop Slave Containers

Start two slave nodes for HDFS and YARN.

```
docker run -itd --network hadoop --hostname hadoop-slave1 --name hadoop-slave1 -p 8042:8042 liliafxi/spark-hadoop:hv-2.7.2
```

```
docker run -itd --network hadoop --hostname hadoop-slave2 --name hadoop-slave2 -p 8043:8042 liliafxi/spark-hadoop:hv-2.7.2
```

- **Port 8042** → NodeManager

To start the containers if they already exist :

```
docker start hadoop-master  
docker start hadoop-slave1  
docker start hadoop-slave2
```

## 4 Access Hadoop Master Container

```
docker exec -it hadoop-master bash
```

- You are now inside the master container shell.

## 5 Start Hadoop Services

```
./start-hadoop.sh
```

If this script is missing, alternatively use:

```
hdfs namenode -format # only first time  
start-dfs.sh
```

```
start-yarn.sh  
jps
```

- `jps` confirms running daemons: NameNode, DataNode, ResourceManager, NodeManager, SecondaryNameNode.

## 6 Prepare HDFS Input

Create input directory:

```
hdfs dfs -mkdir /input
```

Upload input file:

```
hdfs dfs -put purchases.txt /input
```

Verify uploaded files:

```
hdfs dfs -ls /input
```

Check last few lines of the file:

```
tail purchases.txt
```

## 7 Copy WordCount JAR to Container

From **host machine**, copy the JAR file:

```
docker cp "C:\Users\raouf\OneDrive - esi-sba.dz\Desktop\3CS\Big Data\TP01\untitled\target\wordcount-1.0-SNAPSHOT.jar" hadoop-master:/root/
```

Inside the container, verify:

```
ls /root
```

- You should see `wordcount-1.0-SNAPSHOT.jar`.

## 8 Run WordCount MapReduce Job

```
hadoop jar /root/wordcount-1.0-SNAPSHOT.jar tp1.WordCount /input /output
```

If /output exists from a previous run:

```
hdfs dfs -rm -r /output
```

## 9 Check Output

List output directory:

```
hdfs dfs -ls /output
```

- `_SUCCESS` → job completed successfully
- `part-r-00000` → contains WordCount results

View first 20 lines of output:

```
hdfs dfs -cat /output/part-r-00000 | head -n 20
```

## Cleanup / Rerun

- Hadoop cannot overwrite existing output, so you must delete it before rerunning:

```
hdfs dfs -rm -r /output
```

- To completely reset the lab:

```
stop-yarn.sh  
stop-dfs.sh  
docker rm -f hadoop-master hadoop-slave1 hadoop-slave2  
docker network rm hadoop
```